

Q.Calculate the Customer Life Time Value (CLTV) Using 2 Different methods

1. RFM Method
2. Predictive Modelling

```
In [1]: 1 import pandas as pd
        2 import numpy as np
```

```
In [2]: 1 df = pd.read_csv('customer_purchases.csv')
```

```
In [3]: 1 df
```

Out[3]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	12/1/2010 8:26	2.55	17850.0	United Kingdom
1	536365	71053	WHITE METAL LANTERN	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	12/1/2010 8:26	2.75	17850.0	United Kingdom
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
...
541904	581587	22613	PACK OF 20 SPACEBOY NAPKINS	12	12/9/2011 12:50	0.85	12680.0	France
541905	581587	22899	CHILDREN'S APRON DOLLY GIRL	6	12/9/2011 12:50	2.10	12680.0	France
541906	581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4	12/9/2011 12:50	4.15	12680.0	France
541907	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4	12/9/2011 12:50	4.15	12680.0	France
541908	581587	22138	BAKING SET 9 PIECE RETROSPOT	3	12/9/2011 12:50	4.95	12680.0	France

541909 rows × 8 columns

```
In [4]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   InvoiceNo        541909 non-null object
1   StockCode       541909 non-null object
2   Description      540455 non-null object
3   Quantity        541909 non-null int64
4   InvoiceDate      541909 non-null object
5   UnitPrice       541909 non-null float64
6   CustomerID      406829 non-null float64
7   Country         541909 non-null object
dtypes: float64(2), int64(1), object(5)
memory usage: 33.1+ MB
```

```
In [5]: 1 df.isnull().sum()
```

```
Out[5]: InvoiceNo      0
StockCode      0
Description    1454
Quantity       0
InvoiceDate    0
UnitPrice      0
CustomerID    135080
Country        0
dtype: int64
```

RFM METHOD

- Recency
- Frequency
- Monetary

```
In [6]: 1 df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate']).dt.date
```

```
In [7]: 1 df = df.dropna(subset=['CustomerID']) #dropping null values
```

```
In [8]: 1 df['Total_Sales'] = df['Quantity'] * df['UnitPrice'] # calculating total sales
```

C:\Users\91775\AppData\Local\Temp\ipykernel_800\160927535.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['Total_Sales'] = df['Quantity'] * df['UnitPrice'] # calculating total sales
```

CALCULATING RECENCY

```
In [9]: 1 recency_df = df.groupby(by='CustomerID', as_index=False)['InvoiceDate'].max()
2 recency_df.columns = ['CustomerID', 'LastPurshaceDate']
```

```
In [10]: 1 recent_date=recency_df.LastPurshaceDate.max()
2 print(recent_date)
```

2011-12-09

```
In [11]: 1 recency_df['Recency'] = recency_df['LastPurshaceDate'].apply(lambda x: (recent_date - x).days)
2 recency_df.head()
```

Out[11]:

	CustomerID	LastPurshaceDate	Recency
0	12346.0	2011-01-18	325
1	12347.0	2011-12-07	2
2	12348.0	2011-09-25	75
3	12349.0	2011-11-21	18
4	12350.0	2011-02-02	310

CALCULATING FREQUENCY

```
In [12]: 1 df1= df
2 df1.drop_duplicates(subset=['InvoiceNo', 'CustomerID'], keep="first", inplace=True)
3 frequency_df = df1.groupby(by=['CustomerID'], as_index=False)['InvoiceNo'].count()
4 frequency_df.columns = ['CustomerID', 'Frequency']
5 frequency_df.head()
```

C:\Users\91775\AppData\Local\Temp\ipykernel_800\956064957.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df1.drop_duplicates(subset=['InvoiceNo', 'CustomerID'], keep="first", inplace=True)
```

Out[12]:

	CustomerID	Frequency
0	12346.0	2
1	12347.0	7
2	12348.0	4
3	12349.0	1
4	12350.0	1

CALCULATING MONETARY

```
In [13]: 1 monetary_df = df.groupby(by='CustomerID', as_index=False)['Total_Sales'].sum()
2 monetary_df.columns = ['CustomerID', 'Monetary']
3 monetary_df.head()
```

Out[13]:

	CustomerID	Monetary
0	12346.0	0.00
1	12347.0	163.16
2	12348.0	331.36
3	12349.0	15.00
4	12350.0	25.20

MERGING ALL THE TABLES

```
In [14]: 1 temp_df = recency_df.merge(frequency_df,on='CustomerID')
2 rfm_df = temp_df.merge(monetary_df,on='CustomerID')
3 rfm_df
```

Out[14]:

	CustomerID	LastPurshaceDate	Recency	Frequency	Monetary
0	12346.0	2011-01-18	325	2	0.00
1	12347.0	2011-12-07	2	7	163.16
2	12348.0	2011-09-25	75	4	331.36
3	12349.0	2011-11-21	18	1	15.00
4	12350.0	2011-02-02	310	1	25.20
...
4367	18280.0	2011-03-07	277	1	23.70
4368	18281.0	2011-06-12	180	1	5.04
4369	18282.0	2011-12-02	7	3	36.80
4370	18283.0	2011-12-06	3	16	66.75
4371	18287.0	2011-10-28	42	3	80.40

4372 rows × 5 columns

RANKING

```
In [15]: 1 # (lower is better)
2 rfm_df['RecencyRank'] = (rfm_df['Recency'].rank(ascending=True) - 1) / (len(rfm_df) - 1)
3
4 # (higher is better)
5 rfm_df['FrequencyRank'] = (rfm_df['Frequency'].rank(ascending=False) - 1) / (len(rfm_df) - 1)
6
7 # (higher is better)
8 rfm_df['MonetaryRank'] = (rfm_df['Monetary'].rank(ascending=False) - 1) / (len(rfm_df) - 1)
```

```
In [16]: 1 rfm_df
```

Out[16]:

	CustomerID	LastPurshaceDate	Recency	Frequency	Monetary	RecencyRank	FrequencyRank	MonetaryRank
0	12346.0	2011-01-18	325	2	0.00	96.156486	60.626859	95.950583
1	12347.0	2011-12-07	2	7	163.16	4.598490	18.576985	18.142302
2	12348.0	2011-09-25	75	4	331.36	62.445665	35.758408	9.059712
3	12349.0	2011-11-21	18	1	15.00	27.327843	84.991993	75.623427
4	12350.0	2011-02-02	310	1	25.20	94.577900	84.991993	61.816518
...
4367	18280.0	2011-03-07	277	1	23.70	91.672386	84.991993	62.971860
4368	18281.0	2011-06-12	180	1	5.04	80.256234	84.991993	89.533288
4369	18282.0	2011-12-02	7	3	36.80	12.422787	45.676047	51.681537
4370	18283.0	2011-12-06	3	16	66.75	6.977808	5.159003	36.879433
4371	18287.0	2011-10-28	42	3	80.40	46.282315	45.676047	31.892016

4372 rows × 8 columns

NOW CALCULATING SCORES:

AS THE WEIGTHS ARE NOT POVIDED I WILL USE :

$$0.15 * \text{RecencyRank} + 0.28 * \text{FrequencyRank} + 0.57 * \text{MonetaryRank}$$

```
In [17]: 1 rfm_df['RFM_Score'] = 0.15*rfm_df['RecencyRank'] + 0.28*rfm_df['FrequencyRank'] + 0.57*rfm_df['MonetaryRank']
```

```
In [18]: 1 rfm_df=rfm_df.round(0)
2 rfm_df.head()
```

Out[18]:

	CustomerID	LastPurshaceDate	Recency	Frequency	Monetary	RecencyRank	FrequencyRank	MonetaryRank	RF
0	12346.0	2011-01-18	325	2	0.0	96.0	61.0	96.0	
1	12347.0	2011-12-07	2	7	163.0	5.0	19.0	18.0	
2	12348.0	2011-09-25	75	4	331.0	62.0	36.0	9.0	
3	12349.0	2011-11-21	18	1	15.0	27.0	85.0	76.0	
4	12350.0	2011-02-02	310	1	25.0	95.0	85.0	62.0	

PERFORMING SEGMENTATION:

- 0 - 50 - Low valued customer
- 50 - 75 - Medium valued customer
- 76 - 100 - High valued customer

```
In [19]: 1 rfm_df["Customer_segment"] = np.where(rfm_df['RFM_Score'] > 75 , "High Value Customer", (np.where(rfm_df['RFM_Score'] > 50 , "Medium Value Customer", "Low Value Customer")))
```

```
In [34]: 1 final1 =rfm_df[['CustomerID',"RFM_Score",'Customer_segment']]
        2 final1
```

Out[34]:

	CustomerID	RFM_Score	Customer_segment
0	12346.0	86.0	High Value Customer
1	12347.0	16.0	Low value Customer
2	12348.0	25.0	Low value Customer
3	12349.0	71.0	Medium Value Customer
4	12350.0	73.0	Medium Value Customer
...
4367	18280.0	73.0	Medium Value Customer
4368	18281.0	87.0	High Value Customer
4369	18282.0	44.0	Low value Customer
4370	18283.0	24.0	Low value Customer
4371	18287.0	38.0	Low value Customer

4372 rows × 3 columns

Predictive Modelling

```
In [21]: 1 pip install lifetimes
```

```
In [22]: 1 from lifetimes import BetaGeoFitter, GammaGammaFitter
        2 from lifetimes.utils import summary_data_from_transaction_data
        3 from sklearn.preprocessing import MinMaxScaler
        4
```

```
In [23]: 1 pm = df[['CustomerID','Total_Sales','InvoiceDate']]
```

```
In [24]: 1 lf_tx_data = summary_data_from_transaction_data(pm, 'CustomerID', 'InvoiceDate', monetary_
        2 lf_tx_data.reset_index().head()
```

Out[24]:

	CustomerID	frequency	recency	T	monetary_value
0	12346.0	0.0	0.0	325.0	0.000000
1	12347.0	6.0	365.0	367.0	22.993333
2	12348.0	3.0	283.0	358.0	97.253333
3	12349.0	0.0	0.0	18.0	0.000000
4	12350.0	0.0	0.0	310.0	0.000000

```
In [25]: 1 bgf = BetaGeoFitter(penalizer_coef=0.0)
        2 bgf.fit(lf_tx_data['frequency'], lf_tx_data['recency'], lf_tx_data['T'])
```

Out[25]: <lifetimes.BetaGeoFitter: fitted with 4372 subjects, a: 0.02, alpha: 55.62, b: 0.49, r: 0.84>

```
In [26]: 1 t = 10
2 lf_tx_data['pred_num_txn'] = round(bgf.conditional_expected_number_of_purchases_up_to_time(t), 2)
3 lf_tx_data.sort_values(by='pred_num_txn', ascending=False).head(10).reset_index()
```

```
Out[26]:
```

	CustomerID	frequency	recency	T	monetary_value	pred_num_txn
0	14911.0	145.0	372.0	373.0	45.236483	3.40
1	12748.0	114.0	373.0	373.0	5.457982	2.68
2	17841.0	112.0	372.0	373.0	9.544821	2.63
3	15311.0	90.0	373.0	373.0	43.113222	2.12
4	14606.0	88.0	372.0	373.0	10.123523	2.07
5	13089.0	82.0	367.0	369.0	44.713902	1.95
6	12971.0	71.0	369.0	372.0	51.797042	1.68
7	16422.0	66.0	352.0	369.0	61.474394	1.57
8	14527.0	63.0	371.0	373.0	-0.784444	1.49
9	13408.0	54.0	372.0	373.0	76.712407	1.28

```
In [27]: 1 shortlisted_customers = lf_tx_data[lf_tx_data['frequency']>0]
2 shortlisted_customers = shortlisted_customers[shortlisted_customers['monetary_value'] > 0]
3 shortlisted_customers.head().reset_index()
```

```
Out[27]:
```

	CustomerID	frequency	recency	T	monetary_value	pred_num_txn
0	12347.0	6.0	365.0	367.0	22.993333	0.16
1	12348.0	3.0	283.0	358.0	97.253333	0.09
2	12356.0	2.0	303.0	325.0	25.950000	0.07
3	12358.0	1.0	149.0	150.0	142.800000	0.09
4	12359.0	5.0	324.0	331.0	1.380000	0.15

```
In [28]: 1 ggf = GammaGammaFitter(penalizer_coef = 0)
2 ggf.fit(shortlisted_customers['frequency'],
3         shortlisted_customers['monetary_value'])
```

```
Out[28]: <lifetimes.GammaGammaFitter: fitted with 2681 subjects, p: 0.95, q: 1.59, v: 25.06>
```

```
In [29]: 1 ## PREDICTING average transaction value
2
3 lf_tx_data['pred_txn_value'] = round(ggf.conditional_expected_average_profit(
4     lf_tx_data['frequency'],
5     lf_tx_data['monetary_value']), 2)
6 lf_tx_data.reset_index().head()
```

```
Out[29]:
```

	CustomerID	frequency	recency	T	monetary_value	pred_num_txn	pred_txn_value
0	12346.0	0.0	0.0	325.0	0.000000	0.02	40.35
1	12347.0	6.0	365.0	367.0	22.993333	0.16	24.62
2	12348.0	3.0	283.0	358.0	97.253333	0.09	87.49
3	12349.0	0.0	0.0	18.0	0.000000	0.11	40.35
4	12350.0	0.0	0.0	310.0	0.000000	0.02	40.35

```
In [30]: 1 lf_tx_data['CLV'] = round(ggf.customer_lifetime_value(  
2     bgf,  
3     lf_tx_data['frequency'],  
4     lf_tx_data['recency'],  
5     lf_tx_data['T'],  
6     lf_tx_data['monetary_value'],  
7     time=12,  
8     discount_rate=0.01  
9 ), 2)
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\arraylike.py:402: RuntimeWarning: invalid value encountered in log
result = getattr(ufunc, method)(*inputs, **kwargs)

```
In [53]: 1 final2= lf_tx_data.drop(lf_tx_data.iloc[:, 0:6], axis=1)  
2
```

```
In [55]: 1  
2 final2 = final2.sort_values(by='CLV', ascending=False).head(10).reset_index()  
3
```

```
In [58]: 1 min_v = final2['CLV'].min()  
2 max_v = final2['CLV'].max()
```

```
In [60]: 1 final2['CLV'] =((final2['CLV'] - min_v) / (max_v- min_v)) * 100
```

```
In [61]: 1 final2
```

Out[61]:

	CustomerID	CLV
0	18102.0	100.000000
1	17949.0	53.080220
2	14646.0	34.859091
3	17450.0	34.842560
4	16333.0	23.401344
5	16013.0	17.204289
6	13868.0	10.586562
7	12901.0	2.375153
8	12798.0	2.208514
9	15769.0	0.000000

RESULT OF BOTH RFM AND Predictive modelling

In [62]: 1 final1.merge(final2,on='CustomerID')

Out[62]:

	CustomerID	RFM_Score	Customer_segment	CLV
0	12798.0	16.0	Low value Customer	2.208514
1	12901.0	3.0	Low value Customer	2.375153
2	13868.0	19.0	Low value Customer	10.586562
3	14646.0	0.0	Low value Customer	34.859091
4	15769.0	2.0	Low value Customer	0.000000
5	16013.0	1.0	Low value Customer	17.204289
6	16333.0	3.0	Low value Customer	23.401344
7	17450.0	2.0	Low value Customer	34.842560
8	17949.0	0.0	Low value Customer	53.080220
9	18102.0	0.0	Low value Customer	100.000000

In []: 1